

# WikipediaBase

Chris Perivolaropoulos

Sunday 21 February 2016

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>People</b>	<b>1</b>
<b>3</b>	<b>Functionality</b>	<b>1</b>
3.1	get . . . . .	2
3.2	get-classes . . . . .	4
3.3	get-attributes . . . . .	5
3.4	ort-symbols . . . . .	5
3.5	sort-symbols-named . . . . .	5
<b>4</b>	<b>Getting started</b>	<b>6</b>
<b>5</b>	<b>Architecture</b>	<b>7</b>
5.1	Infobox . . . . .	7
5.2	Infobox tree . . . . .	10
5.3	MetaInfobox . . . . .	10
5.4	Article . . . . .	10
5.5	Fetcher . . . . .	11
5.6	Renderer . . . . .	11
5.7	Caching . . . . .	11
5.8	Logging . . . . .	12
5.9	Utilities . . . . .	12
5.10	Pipeline . . . . .	12
<b>6</b>	<b>Provider/Acquirer model</b>	<b>17</b>
6.1	<b>TODO</b> Example . . . . .	17

<b>7</b>	<b>Testing</b>	<b>20</b>
7.1	Unit testing . . . . .	20
7.2	<b>TODO</b> Examples . . . . .	20
<b>8</b>	<b>Synonyms</b>	<b>20</b>
8.1	Good/Bad synonyms . . . . .	21
8.2	Synonym generation . . . . .	22
<b>9</b>	<b>Backend databases</b>	<b>26</b>
9.1	DBM . . . . .	26
9.2	SQLite . . . . .	26
9.3	Redis . . . . .	26
9.4	Postgres . . . . .	26
<b>10</b>	<b>Data sources</b>	<b>26</b>
10.1	HTML . . . . .	26
10.2	API . . . . .	26
10.3	Dumps / Database . . . . .	26
<b>11</b>	<b>Date parser</b>	<b>27</b>
11.1	Parsing with overlays . . . . .	27
11.2	The implementation . . . . .	28
11.3	The dates example . . . . .	28
11.4	Benchmarks . . . . .	28

## 1 Introduction

WikipediaBase base is a backend to START responsible for providing access to wikipedia related information. The WikipediaBase we refer to is a python rewrite of the now deprecated Ruby WikipediaBase.

## 2 People

The python implementation was initially written by Chris Perivolaropoulos and was eventually handed over to

- Alvaro Morales
- Michael Silver

### 3 Functionality

In WikipediaBase, each (supported) Wikipedia infobox is defined as a class, and each (supported) variable in the infobox is defined as an attribute of that class. All WikipediaBase objects belong by inheritance to the superclass `wikibase-term`, which supports the attributes `IMAGE-DATA`, `SHORT-ARTICLE`, `URL`, `COORDINATES`, `PROPER`, and `NUMBER`.

WikipediaBase commands and their return values use lisp-like encoding. WikipediaBase provides the following operations:

#### 3.1 `get`

Given a class, object name, and typed attribute, return the value as a lisp-readable form. Compare Omnibase's `get` operation.

Valid attribute typecodes are `:code` (for an attribute name as in infobox wiki markup) and `:rendered` (for an attribute name in the rendered form of the infobox).

##### 1. Types

Scripts must return a list of typed values. Valid value typecodes are:

###### (a) `:HTML`

A string suitable for rendering as paragraph-level HTML. The string must be escaped for lisp, meaning double quoted, and with double quotes and backslashes escaped with backslashes. The string is not required to contain any HTML codes. For example:

```
(get "wikipedia-sea" "Black Sea" (:code "AREA"))
=> ((:html "436,402 km2 (168,500 sq mi)"))

(get "wikipedia-president" "Bill Clinton" (:
  code "SUCCESSOR"))
=> ((:html "George W. Bush"))

(get "wikipedia-president" "Bill Clinton" (:
  rendered "Succeeded by"))
=> ((:html "George W. Bush"))
```

###### (b) `:YYYYMMDD`

Parsed dates are represented as numbers, using YYYYMMDD format with negative numbers representing B.C. dates. (Unparsable dates are represented as HTML strings using the :HTML typecode.)

```
(get "wikibase-person" "Barack Obama" (:ID "
  BIRTH-DATE"))
=> ((:yyymmdd 19610804))

(get "wikibase-person" "Julius Caesar" (:ID "
  BIRTH-DATE"))
=> ((:YYYYMMDD -1000713))
```

(c) **:CALCULATED**

Typecode for attributes calculated by WikiBase based on characteristics of the article, e.g., *GENDER* and *NUMBER*. See below under Special Attributes for a complete list of calculated attributes.

(d) **:CODE** Deprecated old synonym for :HTML.

(e) **:STRING** Deprecated old synonym for :HTML.

## 2. Special Attributes

Some attributes are special because they are computed by WikipediaBase rather than being fetched from infoboxes, or rather than being fetched directly. These attributes should be specific to wikibase-term, wikibase-person, and wikipedia-paragraphs.

(a) **SHORT-ARTICLE**, **wikibase-term**

The first paragraph of the article, or if the first paragraph is shorter than 350 characters, then returns the first paragraphs such that the sum of the rendered characters is at least 350.

(b) **URL**, **wikibase-term**

Returns the URL of the article as ((:url URL))

(c) **IMAGE-DATA**, **wikibase-term**

Returns a list of URLs for images in the article content (excludes images that are in the page but outside of the article content). If there are no images, should return an empty list. The "best" image should be the first URL in the list; if there is a picture

at the top of the infobox, this is considered to be the best image, or otherwise the first image that appears anywhere in the article. If there is no caption, the caption value should be omitted, e.g., `((0 "Harimau_Harimau_cover.jpg"))` rather than `((0 "Harimau_Harimau_cover.jpg" ""))`.

(d) **COORDINATES, wikibase-term**

Computed from latitude and longitude attributes given in the article or, if none can be found, the infobox. The value is a list of the latitude and longitude, e.g., `((:coordinates latitude longitude))`

(e) **BIRTH-DATE, wikibase-person**

Fetches from the infobox, or, if it is not found, from the article, or, if it is not found, the category information of the article. Always relies on the first date of birth found, matching one of several supported formats. If this attribute has a value, then the object is considered to be a person with respect to the **GENDER** attribute (see below). The value can be a parsed or unparsed date. Parsed dates are represented as numbers, using `YYYYMMDD` format with negative numbers representing B.C. dates. Unparsed dates are strings.

(f) **DEATH-DATE, wikibase-person**

Fetches similarly to **BIRTH-DATE**. Returns the same value types as **BIRTH-DATE**, except if the person is still alive, throws an error with the reply "Currently alive".

(g) **GENDER, wikibase-person**

Computed from the page content based on heuristics such as the number of times that masculine vs. feminine pronouns appear. Valid values are `:masculine` and `:feminine`.

(h) **NUMBER, wikibase-term**

Computed from the page content based on heuristics such as number of times the page's title appears plural. Valid for all objects. Returns `#t` if many, `#f` if one.

i. **PROPER, wikibase-term**

Computed from the page content based on heuristics such as number of times the page's title appears capitalized when not at the start of a sentence. Valid for all objects. Returns `#t` if proper and `#f` if not.

### 3.2 get-classes

Given an object name, return a list of all classes to which the object belongs, with classes represented as lisp-readable strings. Class names are conventionally given in lower case, but this is not an absolute requirement. E.g.,

```
(get-classes "Cardinal (bird)")
=> ("wikibase-term" "wikipedia-paragraphs" "wikipedia
-taxobox")

(get-classes "Hillary Rodham Clinton")
=> ("wikibase-term" "wikipedia-paragraphs" "wikibase-
person" "wikipedia-officeholder" "wikipedia-person
")
```

### 3.3 get-attributes

Given a class name, return a list of all attributes that the class implements (that is, all variables that the infobox implements), as lisp-readable strings. Also sometimes given is the human-readable rendering of the attribute and/or the value typecode for the attribute. Attribute names are conventionally given in upper case, but this is not an absolute requirement. E.g.,

```
(get-attributes "wikipedia-officeholder" "Barack
Obama")
=> ((:CODE "TERM_END3" :VALUE :YYYYMMDD) ...)
```

### 3.4 ort-symbols

sort-symbols takes any number of symbols and sorts them into subsets by the length of the associated article. E.g.,

```
(sort-symbols "Obama (surname)" "Barack Obama")
=> (("Barack Obama") ("Obama (surname)"))
```

### 3.5 sort-symbols-named

`sort-symbols-named` takes a synonym and any number of symbols and sorts the symbols into subsets; if any symbol name is the same as the synonym, it and its subset are sorted to the front. (This should be a case insensitive match, but is it? And again, what's with the subsets?) E.g.

```
(sort-symbols-named "cake" "Cake (TV series)" "Cake (
  firework)" "Cake (film)" "Cake (drug)"
"Cake" "Cake (band)" "Cake (advertisement)" "The Cake
")
=> (("Cake") ("Cake (band)") ("Cake (advertisement)")
    ("Cake (TV series)"))
    ("The Cake") ("Cake (film)") ("Cake (firework)") ("
    Cake (drug)"))
```

## 4 Getting started

The WikipediaBase implementation that we refer to is written in python. Previous implementations were written in Java and Ruby but the language of choice for the rewrite was python for multiple reasons:

- Python is in the pre-graduate curriculum of MIT computer science department. This will ease the learning curve of new members of Infolab.
- Python is a easy to learn and mature language with a rich and evolving ecosystem. This fact eases the introduction of new maintainers even further.

The entire WikipediaBase resides in a git repository in infolab's github organization page

```
git clone git@github.com:infolab-csail/WikipediaBase
```

WikipediaBase depends on multiple other python packages. Fortunately, python is shipped not only with a great package manager, but also with a mechanism called `virtualenv` that isolates installations of a project's dependencies from the rest of the system, thus avoiding problems like version or namespace collisions. The way this effectively works is that the global python installation is half copied half symlinked to a local directory and the

dependencies are installed only in the local sandbox. To create and activate a python virtualenv:

```
$ virtualenv --no-site-packages py
$ . py/bin/activate
$ which python
/the/local/directory/py/bin/python
```

Now that we can safely install anything we want without breaking any global installation

```
pip install -r requirements.txt
```

We will need some extra stuff for WikipediaBase to work:

- Postgresql
- Redis

The installation process of these packages varies across platforms. Both are databases. Their purpose is for caching repeated computations and for storing ahead-of-time computation like infobox markup name to rendered name maps and synonyms.

## 5 Architecture

### 5.1 Infobox

Infoboxes are tables that are commonly used in wikipedia to provide an overview of the information in an article in a semi structured way. Infoboxes are the main source of information for WikipediaBase.




<div> <div><span></span></div> <div><b>Alonzo Church</b></div> </div> <div>  <div>Alonzo Church (1903–1995)</div> </div>	
<b>Born</b>	June 14, 1903 <span>Washington, D.C., US</span>
<b>Died</b>	August 11, 1995 (aged 92) <span>Hudson, Ohio, US</span>
<b>Residence</b>	<span>United States</span>
<b>Nationality</b>	American
<b>Fields</b>	<span>Mathematics, Logic</span>
<b>Institutions</b>	<span>Princeton University</span> (1929–67) <span>UCLA</span> (1967–95)
<b>Alma mater</b>	<span>Princeton University</span>

Figure 1: An example of an infobox

In mediawiki markup terms an infobox is a markup template with a type that gets rendered into html so that the provided information makes sense in the context that it is provided. For example:

```

{{Infobox scientist
| name           = Gerhard Gentzen
| image          = Gerhard Gentzen.jpg
| image_size     =
| alt            =
| caption        = Gerhard Gentzen in Prague,
                  1945.
| birth_date     = {{Birth date|1909|11|24}}
| birth_place    = [[Greifswald]], [[Germany]]

```

```

| death_date      = {{Death date and age
  |1945|8|4|1909|11|24}}
| death_place     = [[Prague]], [[Czechoslovakia]]
| nationality     = [[Germany|German]]
| fields          = [[Mathematics]]
| workplaces      =
| alma_mater      = [[University of Gottingen]]
| doctoral_advisor = [[Paul Bernays]]
| doctoral_students =
| known_for       =
| awards          =
}}
```

will yield:



Figure 2: An example of an infobox

Infobox types are organized into a fairly wide hierarchy. For example `Template:Infobox Austrian district` is a special case of a `Template:Infobox settlement` and each is rendered differently. For our purposes, and to mirror the markup definition of infoboxes, an infobox  $I$  with attributes  $a_i$  and values  $v_i$  is a set of pairs  $(a_i, v_i)$  together with a infobox type  $t$ . Each attribute  $a_i$  and value  $v_i$  have two forms:

- a rendered form,  $a_i^r$  and  $v_i^r$  respectively, which is the rendered HTML representation and
- a markup form,  $a_i^m$  and  $v_i^m$  which is the mediawiki markup code that corresponds to them.

An article may have more than one infoboxes, for example Bill Clinton article has both Infobox Officeholder and Infobox President infoboxes.

The **Infobox** class is the basic data type for accessing information from the infobox of an article. **Infobox**, as well as **Article**, are what one would use were they to use wikipediabase as a python library. The methods provided by an infobox are:

**types** Because we retrieve an infobox based on a symbol name (ie page name), a single **Infobox** may actually be an interface for multiple infoboxes. There is a separate method, based on this one, for getting types in a format suitable for START.

**Value access** is possible provided either  $a_i^r$  or  $a_i^m$ .

**Rendered keys** are provided using the **MetaInfobox** (see below).

**Infobox export** to python types, namely:

- dict for  $a_i^r \rightarrow v_i^r$  or  $a_i^m \rightarrow v_i^m$
- the entire infobox rendered, or in markup form.

## 5.2 Infobox tree

### 5.3 MetaInfobox

The **MetaInfobox** is a subclass of the **Infobox** that provodes information about the infobox, most importantly a map between markup attributes. Say we have an infobox of type  $I$  which has attributes  $a_1, \dots, a_n$ . Each instance of that infobox  $I$  defines

It is an infobox with all the valid attributes and each value is all the names of all attributes that are equivalent to them. Eg An infobox of type Foo that has valid attributes A, B, C and D and A, B and C are equivalent has a meta infobox that looks something like:

Attribute	Value
A	!!!A!!! !!!B!!! !!!C!!!
B	!!!A!!! !!!B!!! !!!C!!!
C	!!!A!!! !!!B!!! !!!C!!!
D	!!!D!!!

### 5.4 Article

The **Article** data structure is responsible for accessing any resource relevant to the article at large.

## 5.5 Fetcher

The fetcher is an abstraction over the communication of WikipediaBase with the outside world. It is a singleton object that implements a specific interface.

Fetchers are organized in an inheriting hierarchy

**BaseFetcher** The baseclass for fetchers, it will return the symbol instead of trying to resolve it in any way

**Fetcher** contains the core functionality of a a fetcher. It will fetch articles from *wikipedia.org*. It is possible to direct it to a mirror but wikipedia-mirror's runtime performance turned out to be prohibitive.

**CachingFetcher** inherits fetcher and retains it's functionality, only it uses Redis to cache the fetched symbols. It is the default fetcher for wikipedia-base.

**StaticFetcher** is a class that implements the **BaseFetcher** interface but instead of reaching out to some data source for the data the return values are statically defined. It is used most notably by **MetaInfobox** to use the **Infobox** functionality to convey arbitrary information.

By default, markup is fetched from the backend. If `force_live` is set to `True`, the markup will be fetched from live wikipedia.org

When tests are ran on TravisCI, we always want to use live data. We check if Travis is running tests by looking at the `WIKIPEDIABASE_FORCELIVE` env variable.

## 5.6 Renderer

Renderers are singleton classes that are useful for rendering mediawiki markup into HTML. Originally the wikipedia sandbox was used by wikipedia-base for rendering pages because it is slightly faster than the API, but the wikipedia-mirror was really slow at this and wikipedia.org would consider it an abuse of the service and block our IP. For that reason we eventually switched to the API with Redis caching, which works out pretty well because **Renderer** objects end up being used only by **MetaInfobox** which has quite a limited scope, making thus cache misses rarely.

## 5.7 Caching

TODO, check what alvaro did with this

## 5.8 Logging

## 5.9 Utilities

1. General
2. Database utilities

## 5.10 Pipeline

When resolving a query WikipediaBase employs a pipeline of modules to figure out what the best way to respond would be.

1. Frontend

WikipediaBase can be used as a library but it's primary function is as a backend to START. The communication between START and WikipediaBase is carried out over a plaintext telnet connection on port {port} using EDN-like sexpressions. The frontend handles the network connection with START, translates the received queries into calls to knowledgebase and then translate the knowledgebase response into properly formulated sexpressions that it sends back over the telnet connection.

2. Knowledgebase

The knowledgebase is the entry point to the rest of wikipediabase. It uses the Provider/Acquirer pattern to transparently provide the frontend with arbitrary methods. Those methods are responsible for choosing whether we are to resort to classifiers or resolvers (or any other mechanism) for answering the query. Available classifiers and resolvers become accessible to the knowledgebase automatically using their base class.

3. Classifiers

Each classifier is a singleton that implements a heuristic for deducing a set of classes of an object. An object may inhibit zero or more classes. There are a couple classifiers available at the moment. Typically a classifier will only deduce whether an object actually inhibits a specific class or not but that is not necessary.

- (a) Term

The `TermClassifier` simply assigns the `wikipedia-term` class. Wikipediabase only deals with wikipedia related information.

(b) Infobox

The `InfoboxClassifier` assigns to a term the classes of the infobox. For example Bill Clinton's page contains the infobox:

```
{{Infobox president
|name           = Bill Clinton
|image          = 44 Bill Clinton 3x4.jpg{{!}}border
|office         = [[List of Presidents of the United States|42nd]] [[President
|vicepresident  = [[Al Gore]]
|term_start     = January 20, 1993
|term_end       = January 20, 2001
|predecessor    = [[George H. W. Bush]]
|successor      = [[George W. Bush]]
|order1         = 40th and 42nd [[List of Governors of Arkansas|Governor of Ar
|lieutenant1    = [[Winston Bryant]]<br>[[Jim Guy Tucker]]
|term_start1    = January 11, 1983
|term_end1      = December 12, 1992
|predecessor1   = [[Frank D. White]]
|successor1     = [[Jim Guy Tucker]]
|lieutenant2    = [[Joe Purcell]]
|term_start2    = January 9, 1979
|term_end2      = January 19, 1981
|predecessor2   = [[Joe Purcell]] {{small|(Acting)}}
|successor2     = [[Frank D. White]]
|office3        = 50th [[Arkansas Attorney General|Attorney General of Arkansas
|governor3      = [[David Pryor]]<br>[[Joe Purcell]] {{small|(Acting)}}
|term_start3    = January 3, 1977
|term_end3      = January 9, 1979
|predecessor3   = [[Jim Guy Tucker]]
|successor3     = Steve Clark
|birth_name     = William Jefferson Blythe III
|birth_date     = {{birth date and age |1946|8|19}}
|birth_place    = [[Hope, Arkansas|Hope]], [[Arkansas]], [[United States|U.S.]]
|death_date     =
|death_place    =
|party          = [[Democratic Party (United States)|Democratic]]
|spouse         = {{marriage|[[Hillary Clinton|Hillary Rodham]]|October 11, 19
|relations      = ''See [[Clinton family]]''
|children       = [[Chelsea Clinton|Chelsea]]
|parents        = [[William Jefferson Blythe, Jr.]]<br>[[Virginia Clinton Kell
```

```
|alma_mater      = [[Edmund A. Walsh School of Foreign Service|Georgetown Univer
|religion        = [[Baptists|Baptist]] {{small|(formerly [[Southern Baptist Co
|signature       = Signature of Bill Clinton.svg
|signature_alt   = Cursive signature of Bill Clinton in ink
|website         = {{url|clintonlibrary.gov|Library website}}
}}
```

And therefore gets the class `wikipedia-president`.

(c) Person

`PersonClassifier` assigns the class `wikibase-person` using a few heretics in the order they are described:

i. Category regexes

Use the following regular expressions to match categories of an article.

- `.* person`
- `^\d+ deaths.*`
- `^\d+ births.*`
- `.* actors`
- `.* deities`
- `.* gods`
- `.* goddesses`
- `.* musicians`
- `.* players`
- `.* singers`

ii. Category regex excludes

Exclude the following regexes.

- `\sbased on\s`
- `\sabout\s`
- `lists of\s`
- `animal\s`

iii. Category matches

We know an article refers to a person if the page is in one or more of the following mediawiki categories:

- `american actors`
- `american television actor stubs`
- `american television actors`



- architects
- british mps
- character actors
- computer scientist
- dead people rumoured to be living
- deities
- disappeared people
- fictional characters
- film actors
- living people
- musician stubs
- singer stubs
- star stubs
- united kingdom writer stubs
- united states singer stubs
- writer stubs
- year of birth missing
- year of death missing

For example Leonardo DiCaprio's page has the following categories:

- Leonardo DiCaprio
- 1974 births
- **Living people**
- 20th-century American male actors
- 21st-century American male actors
- American environmentalists
- American film producers
- American male child actors
- American male film actors
- American male soap opera actors
- American male television actors
- American people of German descent
- American people of Italian descent
- American people of Russian descent
- American philanthropists

- Best Actor AACTA Award winners
- Best Actor Academy Award winners
- Best Drama Actor Golden Globe (film) winners
- Best Musical or Comedy Actor Golden Globe (film) winners
- California Democrats
- Film producers from California
- Formula E team owners
- Male actors from Hollywood, California
- Male actors from Palm Springs, California
- Male actors of Italian descent
- People from Echo Park, Los Angeles
- Silver Bear for Best Actor winners

As it is obvious the list of categories is arbitrary and very far from complete. Multiple methods have been considered for fixing this. Some of them are:

- Supervised machine learning methods like SVM using other methods of determining person-ness to create training sets.
- Hand-pick common categories for person articles determined again with the other criteria

#### 4. Resolvers

Resolvers are also singletons but their purpose is to find the value of the requested property.

- (a) Error
- (b) Infobox
- (c) Person
- (d) Section
- (e) Term

If the article matches one or more of the following categories:

#### 5. Lisp types

Lisp type instances are wrappers for python objects or values that are presentable in s-expression form that START can understand. They

are created either from the raw received query and unwrapped to be useful to the pipeline, or by the answer WikipediaBase comes up with and then encoded into a string sent over telnet to START.

## 6 Provider/Acquirer model

WikipediaBase attempts to be modular and extendible. To accomplish this it is often useful to multiplex multiple sources of the same type of data resource. This is particularly useful when accessing heuristic methods like classifier. To promote modularity and to avoid hard dependencies the provider/acquirer model was created:

A **Provider** is an object through which we can access resources that are stored in a key-value fashion. The **Provider** class offers facilities like decorators to make this provision easy. An **Acquirer** has transparent access to the resources of multiple **Provider**s as if they were a single key value store. This pattern is most notably used for the **KnowledgeBase** to provide the **Frontend** with the way of accessing resources.

### 6.1 TODO Example

We demonstrate the pattern with an example: we will embed a small lisp to python

```
from wikipediabase.provider import Provider, Acquirer
    , provide

class EvalContext(Acquirer):
    def __init__(self, closures):
        super(EvalContext, self).__init__(closures)
        self.closures = closures

    def __call__(self, _ctx, expr):
        if isinstance(expr, list):
            # Handle quotes
            if expr[0] is 'quote':
                return expr[1]

            # Call the lambda
            fn = self(_ctx, expr[0])
            return fn(self, *[self(_ctx, e) for e in
                               expr[1:]])
```

```

        if isinstance(expr, basestring) and expr in
            self.resources():
                return self(_ctx, self.resources()[expr])

    return expr

class Lambda(Acquirer):
    def __init__(self, args, expr, env):
        # Get your symbols from all the available
        closures plus an
        # extra for local variables
        super(Lambda, self).__init__([env] + [Symbols
            ()])
        self.args = args
        self.expr = expr

    def __call__(self, _ctx, *args):
        # Add another closure to the list
        arg_provider = Provider();
        for s, v in zip(self.args, args):
            arg_provider.provide(s, v)

        # Build an eval context and run it
        ctx = EvalContext([arg_provider, Provider(
            self.resources())])
        return [ctx(ctx, e) for e in self.expr][-1]

class Symbols(Provider):
    @provide('setq')
    def setq(self, ctx, symbol, val):
        self.provide(symbol, val)

class Builtins(Provider):
    @provide('lambda')
    def _lambda(self, ctx, args, *body):
        return Lambda(args, list(body), Provider(ctx.
            resources()))

    @provide('if')
    def _if(self, ctx, proposition, then, _else):
        if ctx(ctx, proposition):
            return ctx(ctx, then)
        else:

```

```

        return ctx(ctx, _else)

GLOBAL_EVAL = EvalContext([Builtins(), Symbols()])

```

This little lisp supports:

- lambdas
- A global symbol table
- lexical scoping
- conditionals

It really is very far from being remotely close to a usable language but it can do some cute tricks:

We can evaluate python types:

```

>>> GLOBAL_EVAL({}, 1)
1
>>> GLOBAL_EVAL({}, True)
True
>>> GLOBAL_EVAL({}, "hello")
'hello'
>>> GLOBAL_EVAL({}, list)
<type 'list'>

```

We can define lambdas and call them. The following is equivalent to  $(\lambda a.a)1$ , which should evaluate to 1:

```

>>> GLOBAL_EVAL({}, [{"lambda", ['quote', ['a']], 2,
                        'a'], 1])
1

>>> GLOBAL_EVAL({}, [['lambda', ['quote', []], ['setq',
                        'b', 2], 'b']])
2

```

## 7 Testing

### 7.1 Unit testing

The good functioning of WikipediaBase is assured by a comprehensive test suite of unit tests, functional tests and regression tests.

1. Unit tests

Unit tests test small blocks of functionality, that are composed to create the system at large. For unit testing we use python's default testing library. Each test is a class the subclasses

2. Functional and regression tests

Functional tests are tests written before, during or shortly after the development of a system and they assert the correct overall functioning of the system. Regression tests are very akin to functional tests. They prove that a found bug was fixed and assert that it will not appear again later. Functional and regression tests currently reside in `tests/examples.py`

### 7.2 TODO Examples

## 8 Synonyms

Before we talk about synonyms it is important to concretely define symbols in the context of the omnibase universe:

Symbols are identifiers of "objects" in a data source. (The term "symbol" is unfortunate, since it has so many meanings in computer science, but we're stuck with it for historical reasons.)

Since language tends to have multiple ways of referring to the same things, defining aliases for symbols is imperative.

Synonyms are names which users can use to refer to symbols. (The term "synonym" is unfortunate, because this is really a one-way mapping - "gloss" would be a better term but we're stuck with "synonym" for hysterical raisins.)

The definition of synonyms is the job of the backend itself. Therefore it is the job of WikipediaBase to define the set of synonyms required.

## 8.1 Good/Bad synonyms

There are rules to what is considered a good and what a bad synonym. In short synonyms:

- Should not lead with articles ("the", "a", "an")
- Should not lead with "File:" or "TimedText:".
- Should not fragment anchors. Eg "AlexanderPushkin#Legacy"
- Should not start with the following:
  - "List of "
  - "Lists of "
  - "Wikipedia: "
  - "Category: "
  - ":Category: "
  - "User: "
  - "Image: "
  - "Media: "
  - "Arbitration in location"
  - "Communications in location"
  - "Constitutional history of location"
  - "Economy of location"
  - "Demographics of location"
  - "Foreign relations of location"
  - "Geography of location"
  - "History of location"
  - "Military of location"
  - "Politics of location"
  - "Transport in location"
  - "Outline of topic"
- Should not match `\d\d\d\d in location` or `location in \d\d\d\d`

- Should not be names of disambiguation pages. To make this inclusive for all relevant pages, including typos, that means symbols that match `\([Dd]isambig[^\)]*\)`
- Synonyms that both a) could be mistaken for ones that start with articles and b) might subsume something useful. That means that for example "A. House" (synonym of "Abraham House") is disqualified because it might mislead START in the case of questions like "How much does a house cost in the Silicon Valley?". On the other hand "a priori" can be kept because there are no sensible queries where "a" is an article before "priori".

## 8.2 Synonym generation

To accommodate these restrictions two methods are employed. Disqualification and modification of synonym candidates. First modification is attempted and if that fails we disqualify. The rules for modification are as follows:

- Strip determiners (articles) that are at the beginning of a synonym (or would be at the beginning if not for punctuation):
  - "A "
  - "An "
  - "The "
  - '(The) '
  - The
  - etc.
- Generate both versions, with and without paren. Eg given symbol "Raven (journal)" generate both:
  - "Raven (journal)"
  - "Raven"
- Generate before and after slash, but not the original symbol, e.g.:
  - Given symbol "Russian language/Russian alphabet" generate both
    - \* "Russian language"
    - \* "Russian alphabet"



- Reverse inverted synonyms with commas. Eg given synonym "Congo, Democratic Republic Of The" invert it to get "Democratic Republic Of The Congo"
- As usual, get rid of leading articles if necessary. Eg given synonym "Golden ratio, the" replace it with "the Golden ratio", then strip articles to get: "Golden ratio" same goes for a, an, etc.

This way we generate an initial set of synonyms from the name of the object itself. Furthermore we can generate a set of synonyms from wikipedia redirects to the article. Wikipedia kindly provides an SQL dump for all redirects.

To load the table, in your database where you have loaded the wikipedia data, you should load the redirects table:

```
wget https://dumps.wikimedia.org/enwiki/latest/enwiki
-latest-redirect.sql.gz \
-O redirect.sql.gz && gzcat redirect.sql.gz | mysql
```

And then from the SQL db to find all (good and bad) synonyms to Bill Clinton you can:

```
mysql> select page_title, rd_title from redirect join
       page on rd_from = page_id and (rd_title = "
       Bill_Clinton" or page_title = "Bill_Clinton");
+--
+-----+-----+
| page_title                | rd_title                |
|                             |                           |
+--
+-----+-----+
| BillClinton               | Bill_Clinton            |
|                             |                           |
| William_Jefferson_Clinton | Bill_Clinton            |
|                             |                           |
| President_Clinton         | Bill_Clinton            |
|                             |                           |
| William_Jefferson_Blythe_IV | Bill_Clinton            |
|                             |                           |
| Bill_Blythe_IV            | Bill_Clinton            |
|                             |                           |
```

Clinton_Gore_Administration	Bill_Clinton
Buddy_(Clinton's_dog)	_Bill_Clinton_
_Bill_clinton	_Bill_Clinton_
_William_Jefferson_Blythe_III	_Bill_Clinton_
_President_Bill_Clinton	_Bill_Clinton_
_Bull_Clinton	_Bill_Clinton_
_Clinton,_Bill	_Bill_Clinton_
_William_clinton	_Bill_Clinton_
_42nd_President_of_the_United_States	_Bill_Clinton_
_Bill_Jefferson_Clinton	_Bill_Clinton_
_William_J._Clinton	_Bill_Clinton_
_Billl_Clinton	_Bill_Clinton_
_Bill_Clinton\	_Bill_Clinton_
_Bill_Clinton's_Post_Presidency	Bill_Clinton
Bill_Clinton's_Post-Presidency	_Bill_Clinton_
_Klin-ton	_Bill_Clinton_
_Bill_J._Clinton	_Bill_Clinton_
_William_Jefferson_"Bill"_Clinton	_Bill_Clinton_
_William_Blythe_III	_Bill_Clinton_
_William_J._Blythe	_Bill_Clinton_
_William_J._Blythe_III	_Bill_Clinton_
_Bil_Clinton	_Bill_Clinton_
_WilliamJeffersonClinton	_Bill_Clinton_

	_William_J_Clinton_	_Bill_Clinton_
	_Bill_Clinton's_sex_scandals	Bill_Clinton
	Billy_Clinton	Bill_Clinton
	Willam_Jefferson_Blythe_III	Bill_Clinton
	William_"Bill"_Clinton	Bill_Clinton
	Billlll_Clinton	Bill_Clinton
	Bill_Klinton	Bill_Clinton
	William_Clinton	Bill_Clinton
	Willy_Clinton	Bill_Clinton
	William_Jefferson_(Bill)_Clinton	Bill_Clinton
	Bubba_Clinton	Bill_Clinton
	MTV_president	Bill_Clinton
	MTV_President	Bill_Clinton
	The_MTV_President	Bill_Clinton
	Howard_G._Paster	Bill_Clinton
	Clintonesque	Bill_Clinton
	William_Clinton	Bill_Clinton
	William_Jefferson_Clinton	Bill_Clinton
+--		
	-----+-----+	

46 rows in set (11.77 sec)

## 9 Backend databases

Wikipediabase uses primarily a remote data store that implements the mediawiki interface and attempts to deal with the arising performance issues by aggressively caching pages to a backend key-value based database. The interface with the database is abstracted by using a python-style dictionary interface.

### 9.1 DBM

Several dbm implementations are provided by the python standard library. These include:

1. ndbm

### 9.2 SQLite

### 9.3 Redis

### 9.4 Postgres

## 10 Data sources

### 10.1 HTML

The initial approach to getting the data is to retrieve the normal HTML versions of wikipedia articles and using edit pages to retrieve the mediawiki markup. We invariably use the original wikipedia.org site for performance reasons (See wikipedia-mirror runtime performance section).

### 10.2 API

Mediawiki provides an API for all the required functionality.

1. Performance

### 10.3 Dumps / Database

Direct interface with a local database, besides caching using mdb and/or sqlite was not implemented as part of the thesis. However shortly after caching and compile time data pools in redis and postgres were implemented.

## 11 Date parser

Dateparser resides in a separate package called overlay-parse

### 11.1 Parsing with overlays

The concept of an overlay was inspired by emacs overlays. They are objects that specify the behavior of a subset of a text, by assigning properties to it. An overlay over a text  $t$  in our context is tuple of the range within that text, a set of tags that define semantic sets that the said substring is a member of, and arbitrary information (of type  $A$ ) that the underlying text describes. More formally:

$$\begin{aligned} o_i &\in \text{TextRange} \times \text{Set}(\text{Tag}) \times A & \text{numbers} \\ \text{Text} &\rightarrow \{o_1, o_2, \dots, o_n\} \end{aligned}$$

So for example out of the text

*The weather today,  $\overbrace{\text{Tuesday}}^{o_1}$   $\overbrace{21^{st}}^{o_2}$  of  $\overbrace{\text{November}}^{o_3}$   $\overbrace{2016}^{o_4}$ , was sunny.*

We can extract overlays  $\{o_1, \dots, o_4\}$ , so that

$$\begin{aligned} o_1 &= ( \text{r}(\text{"Tuesday"}), & \{\text{DayOfWeek, FullName}\}, & 2) \\ o_2 &= ( \text{r}(\text{"21^{st}"}), & \{\text{DayOfMonth, Numeric}\}, & 21) \\ o_3 &= ( \text{r}(\text{"November"}), & \{\text{Month, FullName}\}, & 11) \\ o_4 &= ( \text{r}(\text{"2016"}), & \{\text{Year, 4digit}\}, & 2016) \end{aligned}$$

Notice how for all overlays of the example we have  $A = \mathbb{N}$ , as we encode day of the week, day of the month, month and year as natural numbers. We encode more precise type information (ie that a day is inherently different than a month) in the tag set.

Once we have a set of overlays we can define overlay sequences as overlays whose ranges are consecutive, that is their and their tag sets match particular patterns. For example we can search for sequences of overlays that match the pattern

$$p = \text{DayOfMonth, Separator}(/), (\text{Month} \wedge \text{Number}), \text{Separator}(/), \text{Year}$$

to match patterns like 22/07/1991, where *Separator(/)* matches only the character "/"

## **11.2 The implementation**

1. Comparison

## **11.3 The dates example**

## **11.4 Benchmarks**