

# Wikipedia Mirror

Chris Perivolaropoulos

Sunday 21 February 2016

## Contents

<b>1</b>	<b>mediawiki stack overview</b>	<b>1</b>
1.1	Elements of the stack . . . . .	2
<b>2</b>	<b>Setting up</b>	<b>4</b>
2.1	Installing mediawiki extensions . . . . .	6
2.2	Loading mediawiki dumps . . . . .	6
<b>3</b>	<b>The xerces bug</b>	<b>7</b>
3.1	The final solution . . . . .	14
3.2	Covering up with spaces . . . . .	14
3.3	The sed command . . . . .	15
<b>4</b>	<b>Tools</b>	<b>16</b>
<b>5</b>	<b>Automation</b>	<b>19</b>
5.1	Makefiles / laziness . . . . .	20
5.2	Bitnami . . . . .	26
<b>6</b>	<b>Performance</b>	<b>27</b>
6.1	Compile time . . . . .	27
6.2	Runtime . . . . .	28
<b>7</b>	<b>Appendix</b>	<b>28</b>
7.1	script sources . . . . .	28

Wikipedia mirror is a system aiming to automate the creation of a local clone of wikipedia containing only the articles - that is not containing users, discussion and edit history. The automated process includes setting up a server, a database and populating that database with the wikipedia articles.

The purpose for this is to provide the option of accessing wikipedia's dataset independently of wikipedia.org.

## 1 mediawiki stack overview

Wikipedia-mirror builds upon the mediawiki stack provided by bitnami. A service that builds the entire server within the confines of a directory. This is useful because we avoided the overhead of dealing with container or VM technologies and we had direct access to the filesystem of the stack while still having bitnami's build system do the tedious job of orchestrating the various components and separating our server from the rest of the system.

The stack is comprised of

- An http server, in our case apache
- The web application runtime, in our case PHP
- A database, in our case MySQL
- The web application itself, in our case mediawiki

All of the above are provided by the the bitnami mediawiki stack. Xampp used to be go-to for that but it is unmaintained so we decided to go with bitnami which works pretty well.

Once the stack is set up properly the wikipedia dump xml is downloaded and then turned into an sql dump with mwdumper. It could be piped directly to MySQL but extracting can take time and things tend to go wrong during the dumping step.

### 1.1 Elements of the stack

We present each of the elements of the stack in more detail below.

#### 1. Apache

As per wikipedia:

The Apache HTTP Server, colloquially called Apache, is the world's most used web server software. Originally based on the NCSA HTTPd server, development of Apache began in early 1995 after work on the NCSA code stalled. Apache played a key role in the initial growth of the World Wide Web, quickly overtaking NCSA HTTPd as the dominant HTTP

server, and has remained most popular since April 1996. In 2009, it became the first web server software to serve more than 100 million websites.

Apache is developed and maintained by an open community of developers under the auspices of the Apache Software Foundation. Most commonly used on a Unix-like system (usually Linux), the software is available for a wide variety of operating systems besides Unix, including eComStation, Microsoft Windows, NetWare, OpenVMS, OS/2, and TPF. Released under the Apache License, Apache is free and open-source software.

It is fair to say that apache is at least one of the most popular web servers on the internet. wikipedia.org itself seems to be using a more complex stack involving varnish, an HTTP accelerator, and nginx, an alternative, also quite popular HTTP server. We arrive at this conclusion by inspecting the headers returned by wikipedia.org. In the `http://www.wikipedia.org` case we are redirected to the secure domain (pay attention to the `Server:` line):

```
$ curl -s -D - http://www.wikipedia.org -o /dev/  
null  
HTTP/1.1 301 TLS Redirect  
Server: Varnish  
[...]
```

And if we directly ask for `https://www.wikipedia.org` nginx seems to be handling our request:

```
$ curl -s -D - https://www.wikipedia.org -o /dev/  
null  
HTTP/1.1 200 OK  
Server: nginx/1.9.4  
[...]
```

However it is beyond the scope of the project to precisely replicate wikipedia's infrastructure. We focus on the functionality. Therefore due to the popularity, familiarity and by virtue of apache being part of the automatically installable bitnami mediawiki stack, we use it as our server.

## 2. PHP

Mediawiki, which is discussed later, is written entirely in PHP, a popular server side, dynamically typed, object oriented scripting language. PHP is essential and is installed along the bitnami mediawiki stack. PHP is popular among web developers partly due to its support for multiple relational database libraries (including PostgreSQL, MySQL, Microsoft SQL Server and SQLite) and it essentially being structured as a template language generating HTML.

## 3. MySQL

Mediawiki can use a number of different SQL database backends:

- **MSSQL:** An SQL database by Microsoft
- **MySQL:** Using the standard PHP library for MySQL.
- **MySQLi:** An extension to the MySQL backend.
- **Oracle:** A propitiatory SQL database by Oracle.
- **SQLite:** An SQL database that is typically accessed as a library rather than over a client-server scheme as is the case with the other options on the list.

Wikipedia provides multiple dump files for SQL tables of secondary importance in MySQL format (eg. page redirects, categories etc) and suggests **mwddumper** which parses the XML dumps of the wikipedia articles into MySQL. That and bitnami providing it as part of its automatically built stack, make MySQL the obvious choice for the wikipedia-mirror stack.

## 4. MediaWiki

Mediawiki is the heart of wikipedia. MediaWiki is a free and open-source wiki application. It was originally developed by the Wikimedia Foundation and runs on many websites, including Wikipedia, Wiktionary and Wikimedia Commons. As mentioned previously, it is written in the PHP programming language and uses a backend database.

The software has more than 800 configuration settings and more than 2,000 extensions available for enabling various features to be added or changed. On Wikipedia alone, more than 1000 automated and semi-automated bots and other tools have been developed to assist in editing. Most of this is not relevant for our purposes. The only extensions

useful for our purposes are `scriunto` and `parserfunctions` and the only useful settings are related to the name of the site, the name of the database etc and are mostly handled by bitnami.

## 2 Setting up

Following are step by step instructions. First, clone the git repo:

```
$ git clone https://github.com/fakedrake/wikipedia-mirror
$ cd wikipedia-mirror
```

At this point in theory one can run `make sql-load-dumps` which will take care of setting up everything needed to load the database dumps into the working SQL database. Of course for that to happen first a couple of steps need to be carried out:

- Download the wikipedia database dumps in XML format.
- Transform them into a format that MySQL understands.
- Set up the bitnami stack that includes a local install of MySQL
- Load the MySQL dumps into MySQL

All of these steps are encoded as part of the dependency hierarchy encoded into makefile targets and are in theory taken care of automatically, effectively yielding a functioning wikipedia mirror. However this process is extremely long and fragile so it is advised that each of these steps be run individually by hand.

First, download and install bitnami. The following command will fetch an executable from the bitnami website and make a local installation of the bitnami stack discussed above:

```
$ make bmw-install
```

Next step is to make sure `maven`, the java is a software project management and comprehension is installed, required to install and setup `mw-dumper` (see below). You can do that by making sure the following succeeds:

```
$ mvn --version
```

Note: if running on Ubuntu 14.04, you may need to install Maven (for Java) using `sudo apt-get install maven`.

Now everything is installed to automatically download Wikipedia's XML dumps and then convert them to SQL using maven. First maven will be downloaded and built. Then the compressed XML dumps will be downloaded from the wikipedia, they will be uncompressed and finally converted to MySQL dumps using `mwddumper`. This is a fairly lengthy process taking 6 to 11 hours on a typical machine:

```
$ make sql-dump-parts
```

After that's done successfully you can load the SQL dumps to the MySQL database.

```
$ make sql-load-parts
```

Finally the

```
$ make mw-extensions
```

## 2.1 Installing mediawiki extensions

For mediawiki to act like wikipedia a number of extensions are required. The installation process of such extensions is not automated or streamline. To automatically manage this complexity a mechanism is provided for declaratively installing extensions. To add support for an extension to wikiplatform one needs to add the following code in `Makefile.mwextensions` (modifying accordingly):

```
MW_EXTENSIONS += newextension
mw-newextension-url = url/to/new/extension/package.
tar.gz
mw-newextension-php = NewExtensionFile.php
mw-newextension-config = '$$phpConfigVariable = "
value";'
```

And wikipedia-mirror will take care of checking if the extension is already installed and if not it will put the right files in the right place and edit the appropriate configuration files. The entry points for managing extensions are (provided that the name of the registered extension is newextension):

```

make mw-print-registered-extensions # Output a list
                                     of the registed extensions
make mw-newextension-enable         # Install and/or
                                     enable the extension
make mw-newextension-reinstall      # Reinstall an
                                     extension
make mw-newextension-disable        # Disable the
                                     extension
make mw-newextension-clean          # Remove the
                                     extension

```

All registered extensions will be installed and enabled when wikipedia-mirror is built.

## 2.2 Loading mediawiki dumps

Wikipedia provides monthly dumps of all its databases. The bulk of the dumps come in XML format and they need to be encoded into MySQL to be loaded into the wikipedia-mirror database. There are more than one ways to do that.

Mediawiki ships with a utility for importing the XML dumps. However its use for importing a full blown wikipedia mirror is discouraged due to performance trade-offs. Instead other tools like mwdumper are recommended that transform the XML dump into MySQL queries that populate the database.

However the recommended tool for translating the XML dumps into MySQL code is mwdumper. Mwdumper is written in java and is shipped separately from mediawiki. Mwdumper can transform data between the following formats:

- XML
- MySQL dump
- SQLite dump
- CSV

For our purpose we are only interested in the XML -> MySQL dump transformation.

### 3 The xerces bug

Probably the greatest challenge while developing wikipedia-mirror was dealing with a bug in `mwddumper` - the tool for converting wikipedia's XML dumps into MySQL dumps - that makes the tool crash on random articles. Since we did not fully grasp the reason that the bug occurs, we only circumvented it by removing the articles that caused the crash, and since this was a big stumbling block to an otherwise fairly straightforward process, we describe our approach in full detail.

So here is exactly what happens: while `make sql-dump-parts` is running the following is encountered:

```
...
376,000 pages (14,460.426/sec), 376,000 revs
(14,460.426/sec)
377,000 pages (14,458.848/sec), 377,000 revs
(14,458.848/sec)
Exception in thread "main" java.lang.
    ArrayIndexOutOfBoundsException: 2048
        at org.apache.xerces.impl.io.UTF8Reader.read(
            Unknown Source)
        at org.apache.xerces.impl.XMLEntityScanner.
            load(Unknown Source)
        at org.apache.xerces.impl.XMLEntityScanner.
            scanContent(Unknown Source)
        at org.apache.xerces.impl.
            XMLDocumentFragmentScannerImpl.scanContent
            (Unknown Source)
        at org.apache.xerces.impl.
            XMLDocumentFragmentScannerImpl$FragmentContentDispatcher
            .dispatch(Unknown Source)
        at org.apache.xerces.impl.
            XMLDocumentFragmentScannerImpl.
            scanDocument(Unknown Source)
        at org.apache.xerces.parsers.
            XML11Configuration.parse(Unknown Source)
        at org.apache.xerces.parsers.
            XML11Configuration.parse(Unknown Source)
        at org.apache.xerces.parsers.XMLParser.parse(
            Unknown Source)
        at org.apache.xerces.parsers.
            AbstractSAXParser.parse(Unknown Source)
        at org.apache.xerces.jaxp.
```



```

        SAXParserImpl$JAXPSAXParser.parse(Unknown
        Source)
    at javax.xml.parsers.SAXParser.parse(
        SAXParser.java:392)
    at javax.xml.parsers.SAXParser.parse(
        SAXParser.java:195)
    at org.mediawiki.importer.XmlDumpReader.
        readDump(XmlDumpReader.java:88)
    at org.mediawiki.dumper.Dumper.main(Dumper.
        java:142)
make: *** [/scratch/cperivol/wikipedia-mirror/drafts/
wikipedia-parts/enwiki-20131202-pages-articles20.
xml-p011125004p013324998.sql] Error 1

```

Inspecting the makefiles and running `make --just-print sql-dump-parts` we find out that the failing command is:

```

$ java -jar /scratch/cperivol/wikipedia-mirror/tools/
mwdumper.jar --format=sql:1.5 /scratch/cperivol/
wikipedia-mirror/drafts/wikipedia-parts/enwiki
-20131202-pages-articles20.xml-
p011125004p013324998.fix.xml > /root/path/
wikipedia-parts//enwiki-20131202-pages-articles20.
xml-p011125004p013324998.sql

```

Fortunately this does not run for too long so we can safely experiment. Here is the `time` output:

```

26.65s user 1.73s system 78% cpu 35.949 total

```

The error seems to be during reading of the XML dump so it is not specific to SQL output. This could be useful for figuring out which article causes the error, removing which will hopefully resolve the error. To find that out we first try exporting to XML:

```

$ java -jar /scratch/cperivol/wikipedia-mirror/tools/
mwdumper.jar --format=xml /scratch/cperivol/
wikipedia-mirror/drafts/wikipedia-parts/enwiki
-20131202-pages-articles20.xml-
p011125004p013324998.fix.xml > /tmp/just-a-copy.
xml

```

As expected the same error as above is yielded. To then look for the last article that `mwdumper` tried to export we print in reverse order the output xml file, finding the last two occurrences of `<title>` with `grep`. We then reverse again to print them in the original order (note that `tac` is like `cat`, only that yields lines in reverse order):

```
$ tac /tmp/just-a-copy.xml | grep "<title>" -m 2 |
tac
<title>The roaring 20s</title>
<title>Cranopsis bocourti</title> # <- This is
the last one
```

This operation finishes quickly despite `/tmp/just-a-copy.xml` being fairly large because `tac` seeks to the end of the file and reads backwards until `grep` finds the 2 occurrences it is looking for and quits. On ext3 the seek operation does not traverse the entire file. Indeed from the `tac` source code:

```
if (lseek (input_fd, file_pos, SEEK_SET) < 0)
    error (0, errno, _("%s:␣seek␣failed"), quotef (
        file));
/* Shift the pending record data right to make room
for the new.
The source and destination regions probably
overlap. */
memmove (G_buffer + read_size, G_buffer,
        saved_record_size);
past_end = G_buffer + read_size + saved_record_size;
/* For non-regex searches, avoid unnecessary
scanning. */
if (sentinel_length)
    match_start = G_buffer + read_size;
else
    match_start = past_end;

if (safe_read (input_fd, G_buffer, read_size) !=
    read_size)
{
    error (0, errno, _("%s:␣read␣error"), quotef (
        file));
    return false;
}
```

Let's save the path of the original xml file in a variable as we will be using

it a lot. So from now on \$ORIGINAL\_XML will be the path of the original xml.

```
$ export ORIGINAL_XML=/scratch/cperivol/wikipedia-  
mirror/drafts/wikipedia-parts/enwiki-20131202-  
pages-articles20.xml-p011125004p013324998.fix.xml
```

First let's see if there is anything strange going on in the xml file:

```
$ grep "<title>Cranopsis_bocourti</title>" -A 200 -B  
100 $ORIGINAL_XML | less
```

| less is to browse and -A 200 -B 100 means *"show 200 lines after and 100 before the matching line"*. Nothing peculiar was found, so we can't really fix the problem in-place, we will try crudely removing the entire article and hope it works (spoiler alert: it does).

We will try to inspect the parents of the `title` of the breaking article. Fortunately the generated xml is indented so we can find the parents based on that. We count 6 spaces of indentation so we will search backwards from there on each level of indentation. The first line we find on each case will be a direct parent of the article.

```
$ for i in {0..6}; do \  
    echo "Level_$i:"; \  
    tac /tmp/just-a-copy.xml | grep "^_\{$i\}<[~/]" -  
    m 1 -n | tac; \  
done
```

```
Level 0:  
17564960:<mediawiki xmlns="http://www.mediawiki.org/  
xml/export-0.3/" xmlns:xsi="http://www.w3.org  
/2001/XMLSchema-instance" xsi:schemaLocation="http  
://www.mediawiki.org/xml/export-0.3/_http://www.  
mediawiki.org/xml/export-0.3.xsd" version="0.3"  
xml:lang="en">
```

```
Level 1:
```

```
Level 2:
```

```
38: <page>
```

```
Level 3:
```

```
Level 4:
```

```
35: <revision>
```

```
Level 5:
```

```
Level 6:
```

```

26:      <text xml:space="preserve">&lt;!-- This
      article was auto-generated by [[User:Polbot]]. --&
      gt;

```

Looks like the xml is just page tag trees thrown in a grand domain called mediawiki. We could have seen that from the java source too but as expensive as this is, it is much faster than dealing with the source of mwdumper.

The easiest way to cut off this article would be awk but that will take ages and we want to optimize and automate this entire process. First let's try just plain comparing the articles:

```

$ cmp /tmp/just-a-copy.xml $ORIGINAL_XML
/tmp/just-a-copy.xml /scratch/cperivol/wikipedia-
mirror/drafts/wikipedia-parts/enwiki-20131202-
pages-articles20.xml-p011125004p013324998.fix.xml
differ: byte 2, line 1

```

That was fast... Let's see what went wrong:

```

$ head $ORIGINAL_XML
<mediawiki xmlns="http://www.mediawiki.org/xml/export
-0.8/" xmlns:xsi="http://www.w3.org/2001/XMLSchema
-instance" xsi:schemaLocation="http://www.
mediawiki.org/xml/export-0.8/□http://www.mediawiki
.org/xml/export-0.8.xsd" version="0.8" xml:lang="
en">
<siteinfo>
  <sitename>Wikipedia</sitename>
  <base>http://en.wikipedia.org/wiki/Main_Page</
base>
  <generator>MediaWiki 1.23wmf4</generator>
  <case>first-letter</case>
  <namespaces>
    <namespace key="-2" case="first-letter">Media</
namespace>
    <namespace key="-1" case="first-letter">Special
</namespace>
    <namespace key="0" case="first-letter" />

$ head /tmp/just-a-copy.xml
<?xml version="1.0" encoding="utf-8" ?>
<mediawiki xmlns="http://www.mediawiki.org/xml/export

```

```

-0.3/" xmlns:xsi="http://www.w3.org/2001/XMLSchema
-instance" xsi:schemaLocation="http://www.
mediawiki.org/xml/export-0.3/_http://www.mediawiki
.org/xml/export-0.3.xsd" version="0.3" xml:lang="
en">
<siteinfo>
  <sitename>Wikipedia</sitename>
  <base>http://en.wikipedia.org/wiki/Main_Page</
    base>
  <generator>MediaWiki 1.23wmf4</generator>
  <case>first-letter</case>
  <namespaces>
    <namespace key="-2">Media</namespace>

```

The attributes of the xml tags are quite different. We count the numbers of lines in `/tmp/just-a-copy.xml` and hope that the corresponding line number in `$ORIGINAL_XML` will be the same line. If that is so we can ignore the contextual xml information and just blank out the problematic article. We will use `wc` which is also quite fast.

```

$ wc -l /tmp/just-a-copy.xml
17564961 /tmp/just-a-copy.xml

```

And the corresponding line in `$ORIGINAL_XML` would be about:

```

$ sed "17564960q;d" $ORIGINAL_XML
[[Willie Jones (American football)|Willie Jones]],

```

Football... nothing to do with frogs. Looks like there is no avoiding some level of parsing.

## 1. Parsing

We will make the following assumptions to avoid the expensive operation of properly parsing the entire document:

- The XML in the original file is valid
- Any XML within the articles is HTML escaped

First off working with lines is slow because user space code needs to look for newlines. Working bytes delegates work to the kernel, speeding things up considerably. So the `dd` is the right tool for the job. So we will first find at which byte is the article I am interested in.

```
$ grep -b "<title>Cranopsis_bocourti</title>" -m
1 $ORIGINAL_XML
1197420547:      <title>Cranopsis  bocourti</title>
```

This may take a little while but you are stuck with it unfortunately. Our strategy is to make two files: `/tmp/original_tail.xml` that will contain all the data *after* the page we want to remove and `/tmp/original_head.xml` that will contain all the data *before* the page we want to remove.

Now we will use `sed` to look for `</page>` after byte 1197420547 which will be point  $x$  we will and dump the contents of `$ORIGINAL_XML` after point  $x$ :

```
$ dd if=$ORIGINAL_XML skip=1197420547 ibs=1 | sed
'0,/<\page>/d' > /tmp/original_tail.xml
```

Great, that worked! `dd` does not copy in reverse so we will need to do something more complex to construct `/tmp/original_head.xml`. Let's say the position where we found the title of the page we want to remove is  $\alpha = 1197420547$  and the point where the page starts is point  $\beta$ . It is fairly safe to assume that  $\beta > \alpha - 1000$  (we can calibrate the constant 1000 if that assumption is wrong, but it turns out that it isn't). This way we only need to search into 1Kb for `<page>`. Effectively instead of copying the bytes in range  $[0, \beta)$  we are concatenating two ranges  $[0, \alpha - 1000] \cup (\alpha - 1000, \beta)$  by making a subshell that will first output the first range and then output  $(\alpha - 1000, \alpha)$  stopping when it finds `<page>`. Here is the one liner:

```
$ (dd count=$((1197420547-1000)) ibs=1 if=
$ORIGINAL_XML; \
dd if=$ORIGINAL_XML count=1000 skip=$((1197420547-1000)) ibs=1 \
| tac | sed '/<page>/,$d' | tac) > /tmp/
original_head.xml
```

### 3.1 The final solution

All the above was used to compose a script that lives in `data/xml-parse.sh` which is utilized by the makefiles to remove all problematic articles.

If mwdumper fails, we identify the article that caused the breakage and remove it using `xml-parse.sh`. Then we rerun mwdumper. We repeat that until mwdumper succeeds. In total the conflicting articles are about 10-15, and are different depending on the dump being used.

### 3.2 Covering up with spaces

From the above exploration of ways for circumventing the issue of the breaking article we omitted a fairly obvious, but thematically different approach: covering up breaking article with spaces. Once we find out the range in which the page resides we can `mmap` precisely in that part of `$ORIGINAL_XML` and then `memset` covering it up with space characters. The actual implementation lives in `data/page_remover.c`, below we present the call to `mmap`:

```
ctx->off = off - pa_off;
ctx->fd = open(fname, O_RDWR, 0x0666);
if (ctx->fd == -1) {
    perror("open");
    return NULL;
}

ctx->size = len;
ctx->data = mmap(0, len+ctx->off, PROT_READ |
    PROT_WRITE,
    MAP_SHARED, ctx->fd, pa_off);
if (ctx->data == MAP_FAILED) {
    perror("mmap");
    return NULL;
}
```

and the `mmemset`:

```
/* You MIGHT want to thread this but I dont think
   it will make
   * much more difference than memset. */
memset(ctx->data + ctx->off, ' ', ctx->size);
```

Surprisingly this did not fix the mwdumper issue, which points to a possible memory leak on the part of xerces but it is beyond the scope of this project to debug and fix 3rd party tools if we have a choice.

### 3.3 The sed command

Above we kind of glazed over the use the `sed` command but it might be interesting to spend some ink on it. Sed is a unix tool found in `coreutils` that according to its man page is a

stream editor for filtering and transforming text.

The basic premise is that the *"pattern space"*, or the input stream which is a normal unix stream coming from a file, a pipe or just `stdin`, is passed through a programmable pipeline. Either the modified pattern space itself is printed or, with the use of the `-n` flag, selected parts of it. Let's look at the use that we have made for `sed` above

Initially we used `sed` to print a specific line in a file:

```
$ sed "17564960q;d"
```

This `sed` program is separated by a semicolon. Sed iterates over the lines of the input stream and runs each of the ; separated commands on them in sequence until one succeeds. The commands here are `17564960q` and `d`. `17564960q` will quit `sed` once line 17564960 is reached. `d` will discard the current line. So `sed` discards lines until it reaches line 17564960 which it prints and quits.

We then used a `sed` command as part of a series of shell commands piped together in order to print all the lines of a stream after a specific pattern (in our case `</page>`).

```
$ sed '0,/<\</page>/d'
```

This time we have only a single `sed` command, `d`. Sed iterates over the lines in the stream, discarding lines in the range of lines 0 to the line that matches `<\</page>`, effectively only printing lines after `</page>`.

Our final use of `sed` is the inverse of the aforementioned one,

```
$ sed '/<page>/,$d'
```

Here `sed` iterates again over all the lines of the stream this time discarding lines in the range between the first line that matches `<page>` until the final line, denoted with a `$`.



## 4 Tools

A number of tools were developed in assisting the process of manipulating and monitoring the process of loading the dumps into the database. They are presented in details below. Since their source code is fairly concise it is presented in the Appendix.

(a) `sql-clear.sh`

`sql-clear.sh` is a small bash script that truncates all tables from a database. Truncating means leaving the MySQL table schemata unaffected and delete all internal data.

(b) `utf8thread.c`

`utf8thread.c` is another low level program that blanks out all invalid utf-8 characters from a file. We used `pthreads` to speed things up.

(c) `webmonitor.py`

`webmonitor.py` is a python script that sets up a web page that shows live data in the form of a histogram about the progress of the database population. `webmonitor.py` serves a static html page and then deeds it the data over websocket. Webmonitor can show any stream of `<epoc date> <float value>` pairs that it receives in its input. As a sample:

```
$ pip install tornado
```

First install the dependencies of the script. That would be `tornado`, an asynchronous web framework supporting websockets. Also use `tornado` to serve the following page:

```
<!DOCTYPE HTML PUBLIC "-//W3C//DTD HTML 4.01//EN" "http://www.w3.org/TR/html4/strict.dtd">
<html>
<head>
<meta http-equiv="Content-Type" content="text/html; charset=utf-8">
<title>DrNinjaBatmans Websockets</title>

<script type="text/javascript" src="http://code.jquery.com/jquery-1.10.1.js"></script>
```

```

<script type="text/javascript" src="http
    ://code.highcharts.com/highcharts.js">
</script>

<script>
var chart; // global
var url = location.hostname + ':' + (
    parseInt(location.port));
var ws = new WebSocket('ws://' + url +
    '/websocket');
ws.onmessage = function(msg) {
    add_point(msg.data);
};

// ws.onclose = function() { alert('
    Connection closed. '); };

var add_point = function(point) {
    var series = chart.series[0],
    shift = series.data.length > %d;

    chart.series[0].addPoint(eval(point)
        , true, shift);
};

$(document).ready(function() {
    chart = new Highcharts.Chart(JSON.
        parse('%s'));
});
</script>

</head>
<body>
    <div id="container" style="width: 800px
        ; height: 400px; margin: 0 auto"></
        div>
</body>
</html>

```

In essence this page expects to read a stream of values from a websocket at `ws://localhost:8888/hostname` – although it is smart enough to change the `localhost:8888` if you are serving this to another location – and plot them in real time using `highcharts.js`.

The attentive reader may notice that the above is not quite HTML but rather a python formatted string. That is for two reasons: first because the chart configuration is handled by python rather than javascript, second because the width of the graph will be calculated at page load time -ie. by python- and the plot needs to be shifted to only show the most recent points.

```
$ for i in {1..100}; do echo $i; sleep 1;
done | \
awk -oL "{print_\$1/100}" | \
python webmonitor.py
```

This will produce, in 1 second intervals, numbers from 1 to 100. Then it normalizes them using `awk` and feeds them to `webmonitor`. After this command executes we can open the browser and then navigate to `localhost:8888`. We utilize this to remotely monitor the total size of data that `mysql` consumes.

(d) `xml-parse.sh`

Simply removing specific articles fixes the xerces error with UTF8. If the articles are alone the error goes away as well. The `xml-parse.sh` script removes the requested article from the xml file.

```
xml-parse.sh <original-xml-file> <
title_of_article_to_remove> [inplace]
```

if `inplace` is the last argument, the `page_remover.c` will be used to cover the article with spaces. This is much faster. Otherwise the page is just omitted and the result is dumped in stdout. After this script finishes you can run:

```
java -jar tools/mwdumper.jar RESULTING_XML --
format=sql:1.5 > SQL_DUMP
```

(e) `page\remover.c`

As previously discussed, the `xerces` library that `mwdumper` depends on fails, seemingly at random, to process certain pages. To address this issue we remove the pages completely and retry. Since this task is fairly straight forward yet performance sensitive we resorted to writing a small low level program in C to address

it, `page_remove.c`. Page remover accepts as input the path of the XML wikipedia dump, the offset of the article and the size of the article. It then uses the `mmap` system call to random-access the data within the file and fill the article with withespace characters. `page_remover.c` is not threaded as the bottleneck is the HDD IO speed.

## 5 Automation

Creating a wikipedia mirror may seem like a straight forward task but it involves many caveats, nuances and repetitive tasks. Multiple methods of automation were employed to carry out the wide variety of tasks involved into the process.

### 5.1 Makefiles / laziness

The most important part of wikipedia-mirror automation is the `make` build system. Below is an outline of the most basic features of make that constitute it an excellent candidate for automating processes like this.

Make is a build system whereby one can declare required files (targets), dependencies for them, and a set of shell commands that will build those targets. Each target is essentially a finite state machine with two states:

- A file that exists and is up to date with its dependencies and
- A file that either doesn't exist or its modification date is older than that of at least one of its dependencies.

And a sequence of shell commands to transition from the latter to the former state.

For example, save the following as `Makefile` in a project that contains the files `foo.c`, `foo.h`, `bar.c` and `bar.h`:

```
foo.o: foo.c foo.h
    gcc foo.c -c -o foo.o

bar.o: bar.c
```

```

gcc bar.c -c -o bar.o

foobar: foo.o bar.o
gcc foo.o bar.o -o foobar

```

this means that to build `foobar` we need `foo.o` and `bar.o`. And to build `foo.o` and `bar.o` we need `foo.c` and `foo.h`, and `bar.c` and `bar.h` respectively. We also provide commands for building `foo.o`, `bar.o` and `foobar`, which are

- `gcc foo.c -c -o foo.o`
- `gcc bar.c -c -o bar.o`
- and `gcc foo.o bar.o -o foobar`

respectively. Notice that there are no rules for the `.c` and `.h` files. That is because `make` should fail if they are not present. So if we run `make foobar`, `make` will check for `foobar`'s existence and modification date. If `foobar` is missing or its modification date is earlier than its dependencies' (ie `foo.o` and `bar.o`) it will be rebuilt. If any dependencies are missing the same logic is applied to that. This way if we build `foobar` once, and then edit `bar.c` and rerun `make foobar`, `make` will recursively deduce that

- `bar.o` is out of date with respect to its dependency `bar.c`
- When `bar.o` is rebuilt it now has a more recent modification date than `foobar` and therefore the latter is out of date with respect to its dependency `bar.o` so it needs to be rebuilt.

This way `make` can infer a near optimal strategy for building each time the minimum amount of required targets.

Now that we made the basic logic of `make` clear let's dive into some basic features that make our life easier.

#### (a) Phony targets

Some tasks do not result in a file and they need to be run every time `make` encounters them in the dependency tree. For this we have the special keyword `.PHONY:`. Here is an example.

```
.PHONY:
clean:
    rm -rf *
```

This tells make that no file named `clean` will emerge from running `rm -rf *`, and also that even if an up-to-date file named `clean` exists, this target is to be run regardless.

It is worth noting that phony dependencies will always render the dependent target out-of-date. For example:

```
.PHONY:
say-hello:
    echo "hello"

test.txt: say-hello
    touch test.txt
```

When `touch test.txt` will be run **every** time we run `make test.txt` simply because make can not be sure that the phony target `say-hello` did not change anything important for `test.txt`. For this reason phony targets are only meant for user facing tasks.

#### (b) Variables

**makefiles** can have variables defined in a variety of ways. Some cases that are being made use of in `wikipedia-mirror` are presented below.

##### i. Recursively expanded variables

```
OBJECTS = foo.o bar.o

show:
    echo $(OBJECTS)
```

Running `make show` will print `foo.o bar.o` to the console. All variables are substituted for their value by wrapping the variable name in parentheses and prefixing the dollar sign (`$`). Makefile variables have no type, reference to a variable is equivalent to simple string substitution, much like it is in unix shell scripting.

Variables defined using a simple equal `=` sign are recursively expanded. This means that after the variable name is

substituted for the variable content a recursive process keeps expanding emergent variables. This can make variable expansion a very powerful tool. For example:

```
library = foo

foo-libs = -lfoo
foo-includes = -I./include/foo

bar-libs = -lbar
bar-includes = -I./include/bar

libs = $($ (library)-libs)
includes = $($ (library)-includes)

waz:
    gcc waz.c $(includes) $(libs)
```

To demonstrate:

```
$ make --just-print
gcc waz.c -I./include/foo -lfoo
```

The expansion that took place step by step are

```
gcc waz.c $(includes) $(libs)
gcc waz.c $($ (library)-includes) $($ (
    library)-libs)
gcc waz.c $(foo-includes) $(foo-libs)
gcc waz.c -I./include/foo -lfoo
```

Notice how variable names were themselves constructed. Variables can also be defined at the command so in this particular example we could easily switch to the **bar** library:

```
$ make --just-print library=bar
gcc waz.c -I./include/bar -lbar
```

## ii. Simple variables

Sometimes it is not desirable for variables to be expanded indefinitely:

```

kurma = the world $(support1)
animal1 = four elephants
animal2 = tortoise
support1 = supported by $(animal1) $(
    support2)
support2 = supported by a $(animal2) $(
    support2)

all:
    echo $(kurma)

```

Here we try to recursively print an infinite message.

```

$ make --just-print
Makefile:5: *** Recursive variable '
    support2' references itself (eventually
    ). Stop.

```

the variable system of make is total, that is to say variable evaluation can be recursive but it needs to terminate. We can circumvent this by using the := assignment operator.

```

kurma = the world $(support1)
animal1 = four elephants
animal2 = tortoise
support1 = supported by $(animal1) $(
    support2)
support2 := supported by a $(animal2) $(
    support2)

all:
    echo $(kurma)

```

And when we run make we get:

```

make --just-print
echo the world supported by four elephants
    supported by a tortoise

```

basically **support2** is removed from scope when the **support2** itself is substituted.

iii. Automatic variables



Makefile also defines some contextual variables that are defined. The automatic variables defined by `gnu make` are the following

- `$@`: The file name of the target of the rule. If the target is an archive member, then `$@` is the name of the archive file. In a pattern rule that has multiple targets (see Introduction to Pattern Rules), `$@` is the name of whichever target caused the rule's recipe to be run.
- `%`: The target member name, when the target is an archive member. See Archives. For example, if the target is `foo.a(bar.o)` then `%` is `bar.o` and `$@` is `foo.a`. `%` is empty when the target is not an archive member.
- `$<`: The name of the first prerequisite. If the target got its recipe from an implicit rule, this will be the first prerequisite added by the implicit rule (see Implicit Rules).
- `$?:` The names of all the prerequisites that are newer than the target, with spaces between them. For prerequisites which are archive members, only the named member is used (see Archives).
- `$^`: The names of all the prerequisites, with spaces between them. For prerequisites which are archive members, only the named member is used (see Archives). A target has only one prerequisite on each other file it depends on, no matter how many times each file is listed as a prerequisite. So if you list a prerequisite more than once for a target, the value of `$^` contains just one copy of the name. This list does not contain any of the order-only prerequisites; for those see the `$|` variable, below.
- `$+`: This is like `$^`, but prerequisites listed more than once are duplicated in the order they were listed in the makefile. This is primarily useful for use in linking commands where it is meaningful to repeat library file names in a particular order.
- `$|`: The names of all the order-only prerequisites, with spaces between them.
- `$*`: The stem with which an implicit rule matches (see How Patterns Match). If the target is `dir/a.foo.b` and the target pattern is `a.%.b` then the stem is `dir/foo`. The stem is useful for constructing names of related files. In a

static pattern rule, the stem is part of the file name that matched the % in the target pattern. In an explicit rule, there is no stem; so `$$` cannot be determined in that way. Instead, if the target name ends with a recognized suffix (see Old-Fashioned Suffix Rules), `$$` is set to the target name minus the suffix. For example, if the target name is `foo.c`, then `$$` is set to `foo`, since `.c` is a suffix. GNU make does this bizarre thing only for compatibility with other implementations of make. You should generally avoid using `$$` except in implicit rules or static pattern rules. If the target name in an explicit rule does not end with a recognized suffix, `$$` is set to the empty string for that rule.

(c) Functions

Functions are similar to variables in that they also expand into strings. The only difference is that they accept parameter variables.

```
greet = "Hello_$(1)_$(from_$(2))"
john-greets = $(call greet,$1,John)

.PHONY:
all:
    @echo $(call john-greets,Chris)
```

And the output here is

```
$ make
Hello Chris (from John)
```

## 5.2 Bitnami

Bitnami is a family of programs that sets up and manages servers stacks. It contains the entire stack installation within a directory making it both modular and portable while avoiding the fuss of dealing with VMs or containers. Bitnami is not open source so there is no way to tell for sure but my best guess is that it manages this by patching the prefix path of MySQL, apache etc binaries with the installation directory.

Bitnami now supports hundreds of stacks, indicatively the most popular are:

- Oclass
- Joomla
- Drupal
- PrestaShop
- MediaWiki
- Moodle
- ownCloud
- Redmine
- Wordpress

## 6 Performance

### 6.1 Compile time

Compile time includes the time it takes for:

- Downloading all the components of a wikipedia server
- The bitnami stack
  - mwdumper
  - mediawiki-extensions
  - Installing and building those components ( $\sim 1$  min)
  - Downloading the wikipedia dumps
  - Pre-processing the dumps ( $\sim 10$  mins)
  - Populating the mysql database ( $\sim 10$  days)

Builds were done on Infolab's Ashmore. The system's specs are quite high end but the bottleneck was the disk IO so less than 1% of the rest of the available resources were used during the MySQL database population. The specifics of the ashmore machine are:

- **CPU:** Xeon E5-1607 3GHz 4-Core 64 bit
- **Main memory:** 64G

- **HDD:** (spinning disk) 500Gb + 2Tb

Since the main bottleneck was the database population -ie MySQL's performance- great effort and experimentation went into fine tuning MySQL but the speedup achieved was negligible so they were not included in the makefiles.

The backend database engine used by MySQL is InnoDB. Some of the optimization methods attempted are:

- Calibrate the `innodb_buffer_pool_size`. While the available memory in ashmore is fairly large, increasing the buffer pool size up to several GB there was no noticeable difference in database population.
- Change `innodb_flush_method` to `O_DSYNC` to avoid using the `fsync` system call. In short the problem with flushing large mapped files with `fsync` is that `fsync` searches for dirty pages in mapped memory pages linearly making it slower and slower as the file gets larger.
- Calibrate the `innodb_io_capacity`. Unsurprisingly the value of this variable was higher than the bandwidth of the HDD.

The only optimization that actually made a difference in database population speed was to edit the MySQL dump to set:

```
SET AUTOCOMMIT = 0; SET FOREIGN_KEY_CHECKS=0;
```

This allowed InnoDB to do more work in the main memory before committing to the disk and also reduced the overall work by trusting that the keys indicating relation to the database actually point somewhere.

## 6.2 Runtime

Runtime of wikipedia mirror turned out to be too slow to be useful and therefore the project was eventually abandoned. Namely for the full wikipedia dump of July 2014 the load time for the Barack Obama, not taking advantage of caching was at the order of ~30s.

## 7 Appendix

### 7.1 script sources

(a) `pageremover.c`

```
/*
 * Copyright 2014 Chris Perivolaropoulos <
 * cperivol@csail.mit.edu>
 *
 * This program is free software: you can
 * redistribute it and/or
 * modify it under the terms of the GNU
 * General Public License as
 * published by the Free Software Foundation,
 * either version 3 of the
 * License, or (at your option) any later
 * version.
 *
 * This program is distributed in the hope
 * that it will be useful, but
 * WITHOUT ANY WARRANTY; without even the
 * implied warranty of
 * MERCHANTABILITY or FITNESS FOR A
 * PARTICULAR PURPOSE.
 *
 * See the GNU General Public License for
 * more details. You should
 * have received a copy of the GNU General
 * Public License along with
 * this program.
 *
 * If not, see <http://www.gnu.org/licenses
 * />.
 *
 * This should fill a range in a file with
 * spaces. This is an in-place
 * operation so it should be pretty fast.
 *
 * Usage: page_remover PATH OFFSET LENGHT
 */

#include <assert.h>
#include <fcntl.h>
```

```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <semaphore.h>
#include <unistd.h>
#include <unistd.h>

#define USAGE_INFO "page_remover_□PATH_□OFFSET_□
    LENGTH"
#define PRINT(ctx, args...) do { sem_wait(&
    ctx->stdio_mutex);    \
                                printf(args);

                                \
                                fflush(stdout)
                                ;
                                \
                                sem_post(&ctx
                                ->
                                stdio_mutex
                                );    \
    } while(0)

typedef struct context {
    int fd;
    size_t size;
    off_t off;
    sem_t stdio_mutex;
    void* data;
} context_t;

context_t* context_init(char* fname, off_t
    off, size_t len)
{
    context_t * ctx = (context_t*)malloc(
        sizeof(context_t));
    off_t pa_off = off & ~(sysconf(
        _SC_PAGE_SIZE) - 1);

    sem_init(&ctx->stdio_mutex, 0 /* Shared.

```

```

        Usually ignored */ , 1);

PRINT(ctx, "Opening %s at %lu (len: %lu)\n",
      fname, off, len);

ctx->off = off - pa_off;
ctx->fd = open(fname, O_RDWR, 0x0666);
if (ctx->fd == -1) {
    perror("open");
    return NULL;
}

ctx->size = len;
ctx->data = mmap(0, len + ctx->off,
                 PROT_READ | PROT_WRITE,
                 MAP_SHARED, ctx->fd,
                 pa_off);
if (ctx->data == MAP_FAILED) {
    perror("mmap");
    return NULL;
}

return ctx;
}

void context_destroy(context_t* ctx)
{
    if (close(ctx->fd) == -1)
        perror("close");

    if (munmap((void*)ctx->data, ctx->size)
        == -1)
        perror("munmap");

    sem_destroy(&ctx->stdio_mutex);
    free(ctx);
}

int main(int argc, char *argv[])
{
    if (argc != 4)
        fprintf(stderr, USAGE_INFO);

    context_t *ctx = context_init(argv[1],
                                   atoi(argv[2]), atoi(argv[3]));

```

```

        /* You MIGHT want to thread this but I
           dont think it will make
           * much more difference than memset. */
        memset(ctx->data + ctx->off, ' ', ctx->
            size);

        context_destroy(ctx);
        return 0;
    }

```

(b) utf8thread.c

```

#include <assert.h>
#include <fcntl.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/mman.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <semaphore.h>
#include <unistd.h>
#include <unistd.h>

sem_t  stdio_mutex;

#define PRINT(args...) do {sem_wait(&
    stdio_mutex);          \
    printf(args);          \
                           \
    fflush(stdout);        \
                           \
    sem_post(&stdio_mutex); \
                           \
    } while(0)

/* #define DEBUG(args...)          PRINT
    (args) */
#define DEBUG(...)

#define DEFAULT_CHAR ' '
#define WORKERS 8

```



```

#define MESSAGE_DENSITY 1000000000

typedef unsigned long long u64;

#define UTF_LC(1) ((0xff >> (8 - (1))) << (8 - (1)))
#define UTF_CHECK(1, c) (((UTF_LC(1) & (c)) == UTF_LC(1)) && (0 == ((c) & (1 << (7-(1))))))

#define UTF_LEN(x) (UTF_CHECK(6, x) ? 6 : \
    UTF_CHECK(5, x) ? 5 : \
    UTF_CHECK(4, x) ? 4 : \
    UTF_CHECK(3, x) ? 3 : \
    UTF_CHECK(2, x) ? 2 : -1)

struct crange {
    u64 start, end;
};

/* Get return the next character after the
   last correct one. */
inline u64 valid_utf8(u64 c)
{
    char i;
    /* Ascii */
    if ((* (char *) c & 0x80) == 0)
        return c+1;

    /* */
    for (i = UTF_LEN(* (char *) c) - 1; i > 0; i--)
    {
        c++;
        if (!UTF_CHECK(1, * (char *) c)) {
            return (u64) NULL;
        }
    }

    return i < 0 ? 0 : c+1;
}

void* fix_range(void* _r)

```

```

{
    struct crange* r = _r;
    u64 tmp, id = r->start;
    long long unsigned count = 0;

    while ((u64)r->start < (u64)r->end) {
        if (count++ % MESSAGE_DENSITY == 0)
            printf ("[worker: 0x%016llx] Done
                    with %lluK.\n", id, count %
                    1024);

        if (!(tmp = valid_utf8(r->start))){
            PRINT("Invalid char 0x%x (next: 0x%
                    x)\n",
                    *(char*)r->start, *(char*)(r
                    ->start+1));
            *((char*)r->start) = DEFAULT_CHAR;
            (r->start)++;
        } else {
            r->start = tmp;
        }
    }

    PRINT ("[worker: 0x%016llx] OUT\n", id);
    return NULL;
}

void run(u64 p, u64 sz)
{
    int n, i;
    u64 wsize;
    pthread_t workers[WORKERS];
    struct crange rngs[WORKERS];

    wsize = sz/WORKERS + 1;
    printf("Base address: 0x%016llx, step
            size: 0x%016llx\n", p, wsize);

    for (i=0; i<WORKERS; i++){
        rngs[i].start = p + wsize*i;
        rngs[i].end = p + wsize*i + wsize;

        PRINT("Spawning worker %d on range [0x
                %016llx, 0x%016llx), %llu bytes...",
                i, rngs[i].start, rngs[i].end,

```

```

        wsize);
    if ((n = pthread_create(workers+i, NULL
        , fix_range, (void*)(rngs+i)))) {
        PRINT("FAIL\n");
        perror("worker");
        return;
    }
    PRINT("OK\n");
}

PRINT ("Wrapping up...\n");
for (i=0; i<WORKERS; i++) {
    PRINT ("Joining worker %d...", i);
    pthread_join(workers[i], NULL);
    PRINT ("OK\n");
    PRINT("Worker %d went through %llu bytes.\n",
        i, (u64)rngs[i].end - (u64)rngs[i].start);
}
}

int main(int argc, char *argv[])
{
    int fd;
    long long int sz, p;
    struct stat buf;

    sem_init(&stdio_mutex, 0 /* Shared.
        Usually ignored */, 1);

    fd = open(argv[1], O_RDWR, 0x0666);
    if (fd == -1) {
        perror("open");
        return 1;
    }

    fstat(fd, &buf);
    sz = buf.st_size;
    printf("File size: 0x%016llx\n", sz);

    p = (u64)mmap (0, buf.st_size, PROT_READ
        | PROT_WRITE , MAP_SHARED, fd, 0);
    if (p == -1) {

```

```

        perror ("mmap");
        return 1;
    }

    run(p, buf.st_size);

    if (close (fd) == -1) {
        perror ("close");
        return 1;
    }

    if (munmap ((void*)p, buf.st_size) == -1)
    {
        perror ("munmap");
        return 1;
    }

    sem_destroy(&stdio_mutex);

    return 0;
}

```

(c) sql-clear.sh

```

#!/bin/bash
MUSER="$1"
MPASS="$2"
MDB="$3"
MYSQL=$4

# Detect paths
AWK=$(which awk)
GREP=$(which grep)

if [ $# -ne 4 ]
then
    echo "Usage: _$0_{MySQL-User-Name}_{
        MySQL-User-Password}_{MySQL-Database
        -Name}_{MySQL_executable_to_use}"
    echo "Drops all tables from a MySQL"
    exit 1
fi

```

```

TABLES=$(($MYSQL -u $MUSER -p$MPASS $MDB -e '
    show tables' | $AWK '{ print $1}' | $GREP
    -v '^Tables' )

for t in $TABLES
do
    echo "Clearing $t table from $MDB
        database..."
    $MYSQL -u $MUSER -p$MPASS $MDB -e "
        truncate $table $t"
done

```

(d) webmonitor.py

```

"""
Just feed pairs of

<epoch date> <float value>

or even just

<float value>

One way to do that would be

    $ <cmd> > stdbuf -oL awk '{print \$1/$$max
    }' | python webmonitor.py

and I will plot them on port 8888. This will
also pipe the input right
out to the output. Strange input will be
ignored and piped this way,
but this needs to be done by awk as well in
the above example.
"""

import sys
import json
import time

from threading import Thread
from collections import deque

import tornado.websocket as websocket

```

```

import tornado.ioloop
import tornado.web

HTML = """
<!DOCTYPE_HTML_PUBLIC "-//W3C//DTD_HTML_
    4.01//EN" "http://www.w3.org/TR/html4/
    strict.dtd">
<html>
  <head>
    <meta http-equiv="Content-Type" content="
      text/html; charset=utf-8">
    <title>DrNinjaBatmans_Websockets</title>

    <script type="text/javascript" src="http
      ://code.jquery.com/jquery-1.10.1.js"></
      script>
    <script type="text/javascript" src="http
      ://code.highcharts.com/highcharts.js"></
      script>

    <script>
      var chart; // global
      var url = location.hostname + ':' + (parseInt
        (location.port));
      var ws = new WebSocket('ws://' + url + '/
        websocket');
      ws.onmessage = function(msg) {
        add_point(msg.data);
      };

      // ws.onclose = function() { alert('
        Connection closed. '); };

      var add_point = function(point) {
        var series = chart.series[0],
        shift = series.data.length > %d;
        chart.series[0].addPoint(eval(point),
          true, shift);
      };

      $(document).ready(function() {
        chart = new Highcharts.Chart(JSON.parse
          ('%s'));
      });
    </script>

```

```

    </head><body><div id="container" style="
        width: 800px; height: 400px; margin: 0
        auto"></div></body></html>
    """

    config = {
        'visible_points': 10,
        'py_chart_opts': { 'chart': { 'renderTo':
            'container',
                                ,
                                defaultSeriesType
                                ': '
                                'spline'
                                },
            'title': { 'text': '
                DrNinjaBatmans_data
            '},
            'xAxis': { 'type': '
                datetime',
                        ,
                        tickPixelInterval
                        ': '150'
                        },
            'yAxis': { 'minPadding
                ': 0.2,
                        'maxPadding
                        ': 0.2,
                        'title': {
                            'text': '
                                Value',
                                ,
                                ,
                                margin
                                ,
                                :
                                ,
                                80}
                        },
            'series': [{ 'name': '
                Data',
                        'data':
                            [[]]}
        ]
    }

```

```

def date_float(s):
    try:
        date, val = s.split()
    except ValueError:
        val = s.strip()
        date = time.time()

    return int(date), float(val)

def send_stdin(fn=date_float):
    for raw in sys.stdin:
        sys.stdout.write(raw)

        # Ignore strange input.
        try:
            jsn = json.dumps(fn(raw))

            buf.append(jsn)

            for w in websockets:
                try:
                    w.write_message(jsn)
                except websocket.
                    WebSocketClosedError:
                        pass

        except:
            pass

    for ws in websockets:
        ws.close()

class StdinSocket(websocket.WebSocketHandler)
:
    def open(self):
        for i in buf:
            self.write_message(i)

        websockets.append(self)

    def close(self):
        websockets.remove(self)

```



```

class MainHandler(tornado.web.RequestHandler)
:
    def get(self):
        self.write(HTML % (int(config['
            visible_points']),
                                json.dumps(config[
                                    'py_chart_opts'
                                ])))

if __name__ == "__main__":
    application = tornado.web.Application([
        (r"/", MainHandler),
        (r'/websocket', StdinSocket),
    ])
    buf = deque(maxlen=int(config['
        visible_points']))
    websockets = []

    config['args'] = []
    for a in sys.argv[1:]:
        if '=' in a:
            k, v = a.split('=', 1)
            config[k] = v
        else:
            config['args'].append(a)

    Thread(target=send_stdin).start()
    application.listen(8888)
    tornado.ioloop.IOLoop.instance().start()

```

(e) xml-parse.sh

```

#!/bin/bash
#
# Simply removing specific articles fixes the
# xerces error with
# UTF8. If the articles are alone the error
# goes away
# aswell. Extremely weird but that's life.
# Fortunately the article is
# just a stub about some toad (Cranopsis

```

```

        bocourti)
#
# xml-parse.sh ORIGINAL_XML
#       TITLE_OF_ARTICLE_TO_REMOVE [inplace]
#
# if 'inplace' is there the c program will be
#   used to cover the article
# with spaces. This is much faster. Should be
#   anyway. Otherwise the
# page is just ommited and the result is
#   dumped in stdout. Helping
# messages are dumped in stderr After this
#   you can run:
#
# java -jar tools/mwdumper.jar RESULTING_XML
#       --format=sql:1.5 > SQL_DUMP

set -e
set -o pipefail

if [[ $# -lt 2 ]]; then
    echo "xml-parse.sh ORIGINAL_XML
        TITLE_OF_ARTICLE_TO_REMOVE [inplace]"
    exit 0
fi

function my_dd {
    coreutils_version=$(dd --version | head
        -1 | cut -d\ -f3 | colrm 2 2 )
    if [[ $coreutils_version -ge 822 ]]; then
        eval "dd iflag=count_bytes iflag=direct
            oflag=seek_bytes ibs=1M $@"
    else
        echo "Your coreutils may be a bit old (
            $coreutils_version). 822 is the one
            cool kids use." >&2
        eval "dd $@ ibs=1"
    fi
}

ORIGINAL_XML=$1

# Dump a part of the file in sdout using dd.
# Usage:

```

```

# file_range <filename> <first_byte> <start/
# end/length>
#
# Length can be negative
function file_range {
    file=$1
    start=$2
    len=$3

    case $len in
        "end") my_dd if=$file skip=$start ||
            exit 1; return 0;;
        "start") my_dd if=$file count=$start ||
            exit 1; return 0;;
        "") echo "len was empty (file: $file,
            start: $start, len $len). Correct
            format <filename> <byte start> <
            length|'start'|'end'>" 1>&2; exit
            1;;
        *) ;;
    esac

    if [[ $len -gt 0 ]]; then
        # Dump to stdout
        my_dd if=$file skip=$start count=$len
            || exit 1
    else
        skip=$(( $start + ($len) ))
        len=$(( - ($len) ))

        if [[ $skip -lt 0 ]]; then
            skip=0
            len=$start
        fi

        # Dump to stdout
        my_dd if=$file skip=$skip count=$len
            || exit 1
    fi
}

function backwards {
    tac -b | rev
}

```

```

function byte_offset {
    grep -b -o -m 1 -F "$1" | cut -d : -f1
}

# Throw everything but the page in stdout
#
# neg_xml_page "Barack Obama"
function neg_xml_page {
    term="<title>$1</title>"
    title_offset=$(cat $ORIGINAL_XML |
        byte_offset "$term")
    echo -e "\n\tMethod:␣$2(blank␣is␣ok)"
        1>&2
    echo -e "\tsearch␣term:␣$term" 1>&2
    echo -e "\tfile:␣$ORIGINAL_XML" 1>&2
    echo -e "\ttitle␣offset:␣$title_offset"
        1>&2

    # Fail the term is invalid
    if [ -z "$title_offset" ]; then
        echo "Found␣'$title_offset'␣Grep-ing␣(
            cat␣␣$ORIGINAL_XML␣|␣grep␣-b␣-m␣1␣-F
            ␣\"$term\"␣|␣cut␣-d:␣-f1)" 1>&2
        exit 1
    fi

    to_page_start=$((($(file_range
        $ORIGINAL_XML $title_offset -1000 |
        backwards | byte_offset "$(echo␣'<page
        >'␣|␣rev)")+7))
    echo -e "\tto␣page␣start␣(relative):␣
        $to_page_start" 1>&2

    file_range $ORIGINAL_XML $title_offset
        end | byte_offset "</page>" >&2
    echo $((($(file_range $ORIGINAL_XML
        $title_offset end | byte_offset "</
        page>")+7)) >&2
    to_page_end=$((($(file_range $ORIGINAL_XML
        $title_offset end | byte_offset "</
        page>")+7)) # len('</page>') == 7
    echo -e "\tto␣page␣end␣(relative):␣
        $to_page_end" 1>&2

```

```

page_start=$(( $title_offset -
               $to_page_start + 1 ))
echo -e "\tpage_start: $page_start" 1>&2

page_end=$(( $title_offset + $to_page_end ) )
echo -e "\tpage_end: $page_end" 1>&2

echo -e "\tbytes_to_copy: $(( $(du -b $ORIGINAL_XML | cut -f1) - $page_start
                               + $page_end ))" 1>&2

echo "Going to copy $page_start bytes"
1>&2
file_range $ORIGINAL_XML $page_start
start
echo "Finished the first half up to
      $page_start, $(( $(du -b $ORIGINAL_XML
                        | cut -f1) - $page_end )) to go"
1>&2
file_range $ORIGINAL_XML $page_end end
echo "Finished the whole thing." 1>&2
}

# Put stdin betwinn mediawiki tags and into
stdout
function mediawiki_xml {
    (head -1 $ORIGINAL_XML; sed -n "/<
      siteinfo>/,/<\/siteinfo>/p;/<\/
      siteinfo>/q" $ORIGINAL_XML ; cat - ;
    tail -1 $ORIGINAL_XML )
}

# 1: XML File
# 2: Article
# 3: Method (leave blank)
# Assert that the file is there and is not
empty
fsize=$(du -b $ORIGINAL_XML | cut -f1)
if [[ 0 -eq $fsize ]]; then
    echo "ERROR: empty xml file $ORIGINAL_XML"
    1>&2
    exit 1
fi

```

```
echo "Will_remove_article_$2'from_file_$1_(  
size:$fsize)" 1>&2  
if ! neg_xml_page "$2" "$3"; then  
    ret=$?  
    echo "XML_parsing_script_failed" 1>&2  
    exit $ret;  
fi
```